Introduction to KL1

Takashi Chikayama

Institute for New Generation Computer Technology
Mita Kokusai Building 21F
4-28 Mita 1-chome, Minatoku, Tokyo JAPAN

Table of Contents

1 Language Overview

KL1 is a programming language designed for ease of writing parallel processing programs. Its design is based on a concurrent logic programming language Guarded Horn Clauses (GHC). The language is used as the common kernel language for parallel inference machines (PIMs), which were developed in the Japanese Fifth Generation Computer Systems (FGCS) project. The operating system of the PIMs and various application programs are described in this language.

KL1 is a language for parallel symbolic processing.

In symbolic processing, complicated data structures are often required. How such data are represented and management of memory area (or disk storage area) they occupy are vital issues. If too much effort goes in such issues, however, less effort can be put into more essential system design. Programming languages for symbolic manipulation are intended to provide automatic mechanisms to help such management. To be more concrete, *symbolic atoms* gives a standard way to uniquely represent entities with unique names; *automatic memory management* mechanism provides automatic allocation and deallocation of standard data structures. In this respect, KL1 has similar features as other programming languages for symbolic processing, such as Lisp.

For parallel processing, the whole computation should be divided into smaller chunks of subcomputation and they should proceed in coordination with synchronization whenever required. For this purpose, there are many language systems that augments parallel execution and synchronization mechanisms (in cooperation with the operating system) to languages originally designed for sequential processing.

This approach, however, has two crucial problems. One is that, as the basic execution principle of the programming language remains sequential, when to apply parallel execution features must be explicitly specified in the source program. This makes running the same program efficiently on different hardware systems with varying parallelism difficult. The situation is worse when developing a program to solve a problem for which good strategies for parallel processing is yet to be known. Frequently, a decent strategy can only be searched on repeated trials and errors. In such cases, rewriting the program in each of such trial steps might also introduce new bugs, making the development process quite effort-demanding.

Another problem is the difficulty in specifying appropriate synchronization. It is not so easy to have the needs of synchronization always in mind during programming. Moreover, once a synchronization bug is introduced, it might appear in different forms for each execution, because of slight differences in low-level scheduling of parallel processes. This sometimes makes debugging a very difficult job.

KL1 took a different approach. It is not a sequential programming language extended with parallel execution features; it is a born concurrent language, in which all the bits of computation may possibly proceed concurrently. This requires frequent synchronization in principle, but making synchronization fully automatic by introducing its data-flow synchronization mechanism, programmers *cannot* introduce simple synchronization bugs into their programs. In KL1, physical parallelism can be specified separately by the feature called *pragma*. Thus, changes in physical parallelism will not change the correctness of programs. The cost of frequent synchronization may be lowered by program analyses by language implementations.

Thanks to these characteristics of the KL1, the amount of effort in research and development of parallel processing software is greatly reduced. By changing only the pragmas and not changing the logic of the programs, various physical parallel execution strategy can be tested without being annoyed by newly introduced bugs. This enables the programmers concentrate on more essential problem of load distribution method.

2 Execution Mechanism of KL1

This chapter describes the outline of the execution mechanism of KL1.

2.1 A Simple Example

As one of the most simple programs of KL1, let us consider a program of *inverter*, which inverts the input, 0 or 1, to the output 1 or 0.

```
not(In, Out):- In = 0 | Out = 1.
not(In, Out):- In = 1 | Out = 0.
```

These two lines mean the following.

- If the first argument of `not`, that is `In`, is 0, make the second argument `Out` be 1.
- If the first argument of `not`, that is `In`, is 1, make the second argument `Out` be 0.

Each of these lines is called a *clause*, which is the basic unit of program definition.

The heading part of each each clause `not(In, Out)` means that the clause is a part of the definition of the predicate `not`, which has two arguments, which are called `In` and `Out` within this clause. This part is called the *head* of the clause. The word *predicate* is used here as KL1 is a logic programming language, but this can be interpreted as *subroutine* here.

Following the head comes a delimiter `:-`, and the part after this delimiter and before the vertical bar `|` is called the *guard*. The guard gives the condition whether this clause can be applied or not.

Between the vertical bar and the full stop `.` is called the *body* of the clause. This part specifies what to do when the clause is applied. Various things can be written in the body, but, in this example, operations to determine the value of the output argument such as `Out = 1` are described.

To run the inverter program given in the previous example above, we have to add a few lines to make it a complete program. A complete program may look like the following.

```
:- module main.

main :- not(1, X), io:outstream([print(X), nl]).

not(In, Out):- In = 0 | Out = 1.
not(In, Out):- In = 1 | Out = 0.
```

The first line declares that this program to be the main program module. The next line beginning with `main` is the main program. It calls the predicate `not` with arguments `1` and `X`. The code `io:outstream([...])` is for printing out the result and we will not go further into it here.

2.2 Program Execution

When the program given in the previous section is run, the predicate `not` is invoked with two arguments `1` and `X` by `not(1, X)`. A predicate name with its arguments written in paratheses this way is called a *goal*. Goals are a unit of execution of KL1 programs.

The number `1` denotes integer value 1, as can be guessed easily. In KL1, names beginning with upper case letters denote variables. Here, a new variable `X` which appears here the first time is passed as the second argument.

The program execution will proceed as follows.

1. The first formal argument given in the definition `In` corresponds to the actual argument `1`. This satisfies the condition given in the guard of the second clause. The guard condition of the first clause is not satisfied. Thus, the second clause is selected.

2. When a clause is selected, its body is executed. In the body, the second argument `Out`, which corresponds to the actual argument `X` in this case, is given the value `0`.

3. As nothing remains to be done, it's the end of the execution.

As the result of this execution process, the value of the variable `X` given as the second argument is determined to be 0, which will be printed out (details omitted here) and the program terminates normally.

Variables in KL1 are quite different from variables in procedural languages such as C. In procedural languages, variables are value holders where different values might be stored as the program execution goes on. Variables in KL1 is more like variables in mathematics; their value is either defined or not defined and, once defined, it will never change.

If the first argument to be passed is altered to 0, what will be printed out becomes 1, as you may expected.

2.3 Unification

Goals with two expressions in both sides of an equation symbol `=`, such as `In = 0` or `Out = 1` in the above inverter example program, are called *unification*. As the equation symbol is used, these goals roughly mean that both sides have the same value, but the details depend on where they appear.

- Unifications appearing in the guard, which specifies the condition of clause application, mean that the equality of the both sides is a part of the condition. As they only test whether they are equal or not, they are sometimes called *passive unification*. In the above example, `In = 0` tests whether `In` is 0 or not.

- Unifications appearing in the body, which specifies action to be taken when the clause is selected, mean that both sides should be made equal somehow. They are sometimes called *active unification*. In the above example, `Out = 1` determines the value of the undefined variable `Out`, thus actively making both sides equal.

2.4 Syntax of Clauses

Clauses that define a program has the following syntax.

```
PredName(ArgList, ...) :- Guard | Body.
```

Each part has the following meanings.

PredName

gives the name of the predicate (or subroutine, if you like) for which this clause gives (a part of) the definition.

ArgList   determines the correspondence of actual arguments given to the predicate and
            the variables written in the clause definition. In KL1, the scope of the variables
            is confined to one clause definition (or one line of interactive input). Variables
            with the same name but in different clause mean distinct variables.

Guard   specifies the condition needed to be satisfied to apply the clause. Any number
            of goals can be written separated by commas and the condition is considered to
            be satisfied when all of them are satisfied. In the guard, only unifications and
            invocations of certain predicates defined in the language can be written.

Body   specifies the action to be taken when the clause is selected. Like the guard, any
            number of goals can be given here and all the goals will be executed when the
            clause is selected. Unlike in guard, user-defined predicates can be invoked from
            the body, in addition to unifications and language-defined predicates.

Predicates in KL1 can be classified into two categories: *builtin predicates* as defined by
the language and *user-defined predicates* as defined in the programs.

Most of the builtin predicates provides primitive operations for various data. Their
behavior is defined as a part of the specification of the language.

A special builtin predicate is a predicate without any argument *true*. When this appears
in the guard, it means a condition which is always true. In the body, it means to do nothing
(no operation).

2.5 Basic Execution Mechanism

On execution of a goal of user-defined predicates, the guard of the clauses defining the
predicate will be tested first. Then one of the clauses whose guard is satisfied is selected,
and goals in its body (including unifications) will be executed (sometime). Repetition of this
process is the execution process of KL1 programs. A clause can be regarded as a rewriting
rule that converts a goal satisfying the guard condition to the goals in the body.

The basic step of KL1 program execution consists of the following.

1. Some goal is taken out of the set of goals to be executed (As there can be multiple
   goals which are completely the same, it is a multi-set, to be more precise).

2. The goal is matched against program clauses. One of those clauses with their heads
   matching the goal and their guard part satisfied is selected.

3. The goal is converted it to the goals given in the body of the selected clause. This step
   is a *reduction* step, as normally goals are converted into simpler goals.

4. The resultant goals are put back to the set of goals to be executed. The body may
   contain only the simplest goal `true`, in which case no goal is put back to the set.

5. The above steps are repeated until no more goals are to be executed.

6. When the set of goals to be executed becomes empty by repetition of such reduction
   steps, the execution terminates.

2.6 Concurrency, Communication and Synchronization

The rewriting steps described above do not necessarily be made sequentially. As rewriting
steps for multiple goals can be simultaneous, the execution can be actually parallel.

Multiple goals can have `shared variables`. When one goal determines the value of a variable, other goals sharing the same variable can use the value to determine which clause to apply for their reduction. This is the communication mechanism of KL1.

Consider the following goals for the inverter program given above.

```
main :- inverter:not(1, X), inverter:not(X, Y), ...
```

In this example, the two goals share the variable `X`.

In KL1, when multiple goals are to be executed, the order of the execution is not strictly determined. By using the priority mechanism (see Section 6.3 [Goal Priority], page 30), the order can be roughly specified, but it gives only non-strict preference.

Assume that the left-hand side goal is executed first. As the first argument is 1, the value of the second argument `X` is determined to be 0. As the right-hand side goal has this `X` as its first argument, clause to be applied is decided using this value 0, and the right goal determines the value of its second argument `Y` to 1. Thus, the execution yields the value 0 for `X` and 1 for `Y`.

What if the order is exchanged? If the right-hand side goal is tried first, as the value of its first argument `X` has not been determined yet, the guard conditions of both of the clauses defining the predicate `not`, namely `In = 0` and `In = 1`, cannot be tested. In such cases, execution of the goal will be postponed until some future time. This behavior is called *suspension*.

As the right-hand side goal cannot proceed further, the left-hand side goal is the only goal to be executed. Executing the left goal, the value of the variable `X` is determined to be 0. This makes the reason of the suspension of the right-hand side goal vanish. Thus, the right goal can resume its execution. This time, as there is no obstacles in testing the guard conditions, execution proceeds and the value of the variable `Y` is determined to be 1. The final results are the same irrespective of the execution order.

This is the mechanism of KL1 for communication and synchronization between goals. To summarize:

- Communication is through shared variables, and
- Synchronization is automatically done when values are tested in guards.

The only mechanism KL1 provides for selective execution (such as if-then) is the clause selection by guard conditions. If the values required for checking the condition are not available yet, the execution will be automatically suspended. Once a variable is given a value, it will retains the same value never changing it afterwards. Thus, in KL1 programs, there is no possibility of making erroneous decision because of unexpected order of execution.

2.7 More on Execution Mechanism

The basic execution mechanism of KL1 is as described above. There are some more details not explained yet.

- What if more than one clause were applicable?

  Which clause will be selected when guards of the two or more clauses are satisfied? The language specification of KL1 does not define which. It is completely up to the implementation. A mechanism to specify priority between clauses is provided (see Section 6.4 [Clause Priority], page 30), but it only gives preference and the implementation might

not obey it. Actually, in non-shared memory systems like Multi-PSI, if some clause can check the guard condition with data available locally, the implementation prefers to use that clause. Thus, in correct KL1 programs, either the guard conditions should be exclusive or, if they are not exclusive, selecting any of the applicable clauses should yield the desired result.

- What if none of the clauses were applicable?

  What will happen if guards of all the clauses for a predicate are found to be false? For example, what if the predicate `not` in the above example is called as `not(3, X)`. In such cases, the execution leads to a *failure*. Failures, like other unexpected anomalies such as arithmetical overflow, are handled as an exception by the exception handling mechanism of KL1.

- What if a active unification tried to give value to a variable, which already had some value?

  If the value the variable already had is the same as one the unification tried to give, nothing will happen. If they are different, it will be a failure of unification, which will be treated as an exception. Although the behavior is well defined, the programming style of unifying two variables both of which already have values is not recommended. Such operation is not only inefficient but also makes the program harder to read.

2.8 Summary of Execution Mechanism

The most essential parts of the language specification of the KL1 has been described, in this section.

- Basic execution mechanism as concurrent reductions,
- Communication between goals by shared variables, and
- Synchronization by suspending guard condition tests.

3 Data Types and Notation

Up to now, only concrete values we used in the examples are integer values. Of course, KL1 programs can handle data of various other types. This chapter describes basic data types provided by KL1 and their notation.

3.1 Logical Variables

One of the most notable difference between logic programming languages such as KL1 and procedural languages such as C is in their *variables*. To clearly distinguish them, variables in logic programming languages are sometimes called *logical variables*.

3.1.1 Variables are not Memory Addresses

In procedural languages, variables are names of some memory area beginning at certain memory address. As they refer to computer memory, their contents can be read out or *overwritten*. Variables in logic programming languages do *not* correspond to memory addresses but they are names of *values*. Thus, their value may be referenced but they are *never overwritten*.

For example, arguments of a C function refer to memory addresses where the arguments are stored. Thus, they can be read out or overwritten. Arguments of a KL1 predicate refer to the value passed to the predicate, *not* the address where the passed value is stored. Thus, they can never be overwritten.

A notable feature of logical variables is that they can be undefined at some moment. Although overwriting is impossible, programs can *define* their value. This is how variables are used to pass values around. When two parts of a program shares a variable, one can define its value and the other can use the value.

3.1.2 Variables Without Types

In KL1 programs, types of values variables can have are not explicitly declared. The same variable in the same clause in the source program (but not the same variable in execution time) can hold values of different types in one invocation and another. This is similar to Prolog or Lisp.

This has both merits and demerits. One demerit is that an efficient implementation becomes more difficult, as what kind of value a variable can have cannot be determined until run time, in principle. This problem can be solved at least partially by static analysis called *type inference* during compilation. For human programmers, the type of the values of variables might have been a great help in understanding the behavior of programs. In this respect, so-called strongly typed languages in which all variables have declared types might be better.

On the other hand, the merits are: no special notation such as *union* in C is needed to use variables which have different kind of data depending on other data; no special mechanism such as *generic package* in Ada is needed to write programs that can be used in common with different types of data.

In one sentence, variables without types make reading programs more difficult, but writing programs easier. This makes the language suitable for experimental programs,

which will be rewritten many times in trial and error style program development process. This might be one of the reasons why Lisp has been used widely in AI research.

3.2 Atomic Data

Atomic data are data without internal structure and the value itself (rather than its contents) has the meaning. KL1 has the following atomic data types.

Symbolic Atom

> Symbolic atoms have no specific meaning. Only their identity is significant. They are useful in uniquely representing various notions in programs.

Integer   Integers represent usual integer values. The range of integers of the KLIC implementation depends on installation; systems with 32-bit words have the range between `-2^27` and `2^27-1` and systems with 64-bit words have the range between `-2^59` and `2^59-1`, inclusive. Normal decimal notation with optional sign is used, such as `3` or `-15`. Bases other than ten can also be specified before an apostroph; `16'1a` means 26. Character codes can be written as `0'a`, which means the code for the lowercase letter `a`.

Floating Point Number

> Floating point numbers approximate a real number. Decimal notation with a decimal point is used as in `3.14`. Exponent part can be added optionally, as in `0.314e+1`.

3.2.1 Symbolic Atoms

Symbolic atoms in KL1, unlike those in Lisp, do not have any attributes. Thus, the only possible operation on symbolic atoms is equality checks in the guard.

The notation of symbolic string is similar to Edinburgh Prolog, which is one of the following.

- A lower case letter followed by a sequence of any number (including zero) of letters, digits or underlines.

  Examples:

  ```
  icot    kl1    a_symbolic_atom_with_a_long_name
  ```

- A sequence of special characters (some of ~, +, -, *, /, \, ^, <, >, =, ' (backquote), :, ., ?, @, #, $, &).

  Examples:

  ```
  +    >=    :-    =:=
  ```

- A sequence of any characters quoted by single quotes. If single quote characters are to be included, they should be doubled.

  Examples:

  ```
  'Hello world'    '''quoted'''
  ```

- Special one character atoms. There are three of them, which are !, | and ;.

- The special atom `[]`, which usually is used to represent ends of lists (see Section 3.3.4 [Lists], page 12).

3.2.2 Numerical Atoms

Arithmetics and comparison on integer and floating-point numbers can be effected by builtin predicates provided as language primitives.

Magnitude comparison on integers can be made by writing builtin predicates in guard conditions. It is much more convenient to use macro notations than to directly use the builtin predicates.

For comparison of integers in the guard, binary operators `=:=`, `=\=`, `<`, `>`, `=<` and `>=` can be used. For example, `X =< Y` tests whether "X is less than or equal to Y". Operators `=:=` and `=\=` can be used for equality and inequality respectively,. Equality can also be tested by unification but, as is also true with all the macros for integer comparison listed here, `=:=` allows arithmetical expressions as its operands.

For example, a predicate that returns the larger of the first and the second arguments to the third argument can be defined as follows.

```
max(X, Y, Z) :- X >= Y | Z = X.
max(X, Y, Z) :- X =< Y | Z = Y.
```

In this program, if `X` and `Y` are the same integer, the guard conditions of the both clauses hold. As mentioned above (see Section 2.7 [More on Execution Mechanism], page 6), the language does not specify which of the two clauses will be selected in such cases. In this program, however, there will be no problem because both yields the same result.

Integer arithmetics are also provided as builtin predicates that can be executed in both the guard and the body, but macro notations are also convenient for them. The variable to be given the value should be written in the left-hand side of a binary operator `:=`, and arithmetical expression using operators such as `+`, `-`, `*`, `/`, `mod`, and parentheses in the right-hand side. See the manual for description of the complete set of operators.

For example, a predicate that returns the sum of the first and the second argument to the third argument can be defined as follows.

```
sum(X, Y, Z) :- true | Z := X + Y.
```

Here, the guard is `true` which means always true, i.e., that this clause can be applied unconditionally.

Almost the same set of operations are available for floating-point numbers, but they are operations different from those for integers. They are distinguished by putting a dollar sign in front of the names of the comparison and assignment operators (but not in front of the arithmetical operators). For example, floating-point versions of the above two example predicates `max` and `sum` can be defined as follows.

```
max(X, Y, Z) :- X $>= Y | Z = X.
max(X, Y, Z) :- X $=< Y | Z = Y.

sum(X, Y, Z) :- true | Z $:= X + Y.
```

3.3 Structured Data

Structured data consists of element data gathered in one way or another, depending on their types. KL1 has the following structured data types. Other data structures such as

one representing executable object code are also available, which will not be discussed in detail here.

Functor   Functors are a data structure with a name and a fixed number of elements. Functors are usually used as a record structure (such as `struct` of C) with their individual elements used for different purposes.

Vector   Vectors are a structure consisting of values of arbitrary data types as its elements. Unlike functors, elements of vectors usually contain the same kind of information but they can be of different types. Each element can be accessed by integer index.

String   Strings are a structure consisting of integer values within some range. Depending on the range of the element values, 1-, 8-, 16- and 32-bit strings are available. Strings are basically the same as vectors but, as the type of the elements and their values are restricted, more compact representation and more efficient operations are possible.

List   Lists actually are not basic data type provided by the language. A list consists of zero or more of data structured named *cons*. Cons is a structure with two elements. One cons is essentially nothing more than a vector with two elements, but by making a combination of multiple conses, more flexible list structures can be represented.

3.3.1 Functors

Functor structures are structures with given name (*principal functor*) and one or more elements, which can be of any type.

Functor constants can be written by the name of the principal functor, a left parenthesis, elements separated by commas, and finally a right parenthesis. Functor names have the same syntax as symbolic atoms. The principal functor name and the following left parenthesis should *not* be separated by space characters or any other punctuation symbols. Elements can be of any type, including variables or functors themselves.

Examples:

```
f(a, 3)   'a recursive functor structure'(X, 'child functor'(Y))
```

Functors are conveniently used for representing data structures whose size is known beforehand; they correspond to record structures such as `struct` of C. Only unification (both passive and active; see Section 3.3.6 [Unification of Structures], page 13) are normally used for functor structures except in metaprograms (programs manipulating programs).

3.3.2 Vectors

Vectors are a data structure which can have any number of (possibly zero) elements of arbitrary types.

Vector constants are written by an open brace `{`, elements separated by commas, and then a close brace `}`. Elements can be of any type, including variables. A vector with no elements is written as `{}`.

Examples:

```
{}   {1, 2, 3}   {a, b, f(c), X, {Y, d}}
```

Elements of vectors are indexed from zero. For example, the element with index 1 of the vector `{a, b, c}` is `b`.

3.3.3 Strings

A string is a data structure which has integer values within certain range as its elements. It is often used to represent character strings which has codes character as its elements. Such character strings can be written as, for example, `"abc"`, writing the element characters within double-quote symbols. In the KLIC implementation, the default range of the elements is 0 through 255 (unsigned 8 bit integers) and ASCII character code set is used. Characters of 16 bit code sets can also be used by EUC (Extended Unix Code) coding scheme, that uses two elements of strings for one character.

Elements of strings are also indexed by integers beginning from zero. For example, the element with index 1 of the string `"abc"` is the character code for `b`.

3.3.4 Lists

As in Lisp, lists of KL1 are constructed using *cons* data structures.

A cons is a data structure with two elements. Following Lisp, these elements are called *car* and *cdr*, respectively. For their notation, two elements are written within brackets `[` and `]`, separated by a vertical bar `|`. For example, `[a|b]` denotes a cons whose car is a symbolic atom `a` and cdr `b`.

When cons structures are used to form a list structure, the car part of a cons represents one element of the list and the cdr part represents the rest. A list without any elements is represented as the symbolic atom `[]`. A list with two elements will be `[a | [b | []]]` using nested conses. As this notation is not quite convenient for longer lists, this can also be written as `[a, b]`. In general, lists can be written by an open bracket, elements separated by commas, and finally a close bracket.

As lists are data structures used very frequently in KL1, and understanding their syntax is crucial in understanding what follows, let us explain some more examples.

The basic notation for lists is `[Car|Cdr]`, which consists of the first element *Car* and the tail of the list *Cdr*. An empty list is represented by an atom `[]`. If *Cdr* happens to be empty, that is, when the list consists of only one element *Car*, such a list can be written as `[Car|[]]` in the basic notation, or, alternatively, as `[Car]`.

A list consisting of four elements, `one`, `two`, `three` and `four` can be written as `[first, second, third, fourth]`. On the other hand, a list consisting of four or more elements, but with the first four elements being `first`, `second`, `third` and `fourth`, can be written as `[first, second, third, fourth | Rest]`. Here, the variable `Rest` corresponds to the list beginning with the fifth argument, or an empty list if the whole list had only four elements.

3.3.5 Incomplete Data Structures

Data structures with elements of arbitrary types, such as vectors and lists, can have values of their elements left undetermined. Such data structures are called an *incomplete data structure*.

Elements of incomplete data structures which do not have their value determined yet are written as variables. For example, `[a|X]` means an incomplete list whose first element is a symbolic atom `a` and the rest of the list is represented as a variable `X`.

Variables in incomplete data structures can also be shared among goals. Their values can be determined by goals other than the goal which has the whole data structure. Thus, incomplete structures can be gradually made more complete by some other goals. Incomplete data structures play important roles in various programming technique of KL1, some of which will be described later in this document. They are the source of flexible programming styles of KL1.

3.3.6 Unification of Structures

Structured data also can be operands of unification.

3.3.6.1 Guard Unification of Structures

As stated above, guard unifications check out whether the both operands have the same value or not. For structured data, the unification rule becomes recursive.

1. First, check if both operands have the same type and the same number of elements. Otherwise, they are not equal.

2. Then, apply unification to all the corresponding elements (with the same index values) of two structures. When and only when all the elements are equal, the two structures are equal as a whole.

On unification of elements, if the element of one of the operands is a new variable, the corresponding element of the other operand is given to the variable as its value. For example, the following program defines a predicate which, when a list is given as its first argument, returns its first element to the second argument, and the tail of the list to the third argument.

```
carcdr(Cons, Car, Cdr) :- Cons = [X|Y] | Car = X, Cdr = Y.
```

Note that, when the first argument is not a cons, the unification in the guard will not succeed and thus this clause cannot be applied, according to the first rule of the guard unification given above.

3.3.6.2 Body Unification of Structures

When one of the operands of a body unification is a variable whose value is not determined yet, its value is determined to be the other operand, whether it is structured or atomic.

If both operands are structures, the unification rule is recursive as follows.

1. First, check if both operands have the same type and the same number of elements. Otherwise, the unification fails.

2. Then, apply unification to all the corresponding elements (with the same index values) of two structures. When and only when all the elements are successfully unified, the two structures are successfully unified.

Note however, although the above are the specification, applying body unification to two operands both having their values already determined is not a recommended coding style.

3.3.7 Structures are Values

It may be worth noting that structures in KL1 are structured values, unlike structured storage in procedural languages that has ever changing values. If two structures are the same at one time, they will be the same forever after in KL1.

In procedural languages that allow assignments to structure elements, *identity* of structures is an important notion. Even if two structures have the same value at one time, one of them may be modified later making them different.

In logic programming languages, there is no distinction between *equality* and *identity* as only values are of concern. Once two structures are unified, that is, their equality is guranteed in a guard unification or they are made equal in a body unification, they both can be considered to be identical thereafter, as they will never change the values of their elements.

Although incomplete data structures introduce some subtlety, lack of changes of structure elements makes distributed memory implementation of the language much easier. Each processor can keep its own copy of shared data structures, rather than accessing their origin. This allows much more efficient processing.

3.4 Notational Conventions

Before going into programming technique issues, let us introduce several notational conventions.

3.4.1 Guard Unification within Head

Guard unification with one of the arguments passed to the predicate can be more tersely written by writing the pattern to match against it directly at the position of the argument in the head. The following two clauses have the same meaning.

```
not(In, Out) :- In = 1 | Out = 0.
not(1, Out) :- true | Out = 0.
```

Unification may be against a structure and that structure may have variables in them. The structure to be matched against can also be written in the head of clauses. For example, the following two clauses have the same meaning.

```
carcdr(Cons, Car, Cdr) :- Cons = [X|Y] | Car = X, Cdr = Y.
carcdr([X|Y], Car, Cdr) :- true | Car = X, Cdr = Y.
```

A similar convention applies to when the guard unification is not against a constant but with another variable which is an argument. The following two clauses, that check the equality of two arguments, have the same meaning.

```
same(X, Y, Ans) :- X = Y | Ans = same.
same(X, X, Ans) :- true | Ans = same.
```

3.4.2 Omitting the Goal true

When the guard of a clause is unconditional, i.e., when the guard has invocations of the predicate `true` only, the whole guard can be omitted along with the vertical bar separating the guard and the body. For example, the following two clauses have the same meaning.

```
not(1, Out) :- true | Out = 0.
not(1, Out) :- Out = 0.
```

When not only the guard but also the body had only the invocations of `true`, the body also can be omitted along with symbol `:-` that separates the head and the body. For example, the following two clauses means the same.

```
one(1) :- true.
```

```
one(1).
```

3.5 Summary of Data Types and Notation

Data structures and operations on them available in KL1 are described in this chapter.

- Variables of KL1 are logical variable that are names of values rather than names of memory storage.
- KL1 provides symbolic atoms and numerical atoms. Numerical atoms include integers and floating point numbers.
- As structured data, KL1 provides functors, vectors, strings and lists. A characteristic feature of KL1 is incomplete structures that are data structures with some of their elements undetermined.

4 Processes and Streams

Basic features of the KL1 language were described in the previous sections. In this section, a programming technique very frequently used in KL1 is described, in which processes and streams connecting them are used.

4.1 Processes

As explained above, the basic execution mechanism of KL1 is concurrent (that is, possibly parallel without specific order) repetition of reductions of goals. This bare mechanism, however, is not in sufficiently high level for describing complicated computation in a way easily understood. The notion of *process* helps grouping the repeated reductions making higher level computation structures.

4.1.1 Processes of KL1

Predicates with an invocation of the same predicate (probably with slightly different arguments) in their body can be regarded as expressing an execution loop. One reduction of a goal of such a predicate will result in another goal for the same predicate, which is to be reduced again. For example, a predicate that counts down to zero starting with the given argument can be defined by the following two clauses.

```
count_down(0).
count_down(N) :- N > 0 | M := N-1, count_down(M).
```

In this example, counting down is *not* done by `N := N-1`. As stated above, variables in KL1 are not places to store changing values; rather, they are names given to some values themselves. Thus, a new name `M` is needed for `N-1`, which is a value different from `N`.

The body of the second clause in the above example contains two goals: `M := N-1` that performs subtraction and `count_down(M)`, recursive invocation. As the order of goals in the body does not make any significance in KL1, their execution order is not defined. The recursion may precede the subtraction and they may also be executed in parallel. As both of the clauses for the predicate `count_down` have a guard condition or its abbreviation in the head that tests the argument, the reduction of the recursive goal will anyway be suspended until the value of the argument, i.e., the result of subtraction, is obtained.

Although each reduction is only a fragment of computation, the repetitive reductions of recursive invocations as a whole can be regarded as a computation process of considerable size. This repetitive reductions is called a *process*. Note that, in KL1, the notion of processes is *not* a part of the language, but a notion available in a programming technique, although the technique is so commonly used that it virtually is a part of the language. In the rest of this section, the characteristics of processes in KL1 are explained with concrete examples.

4.1.2 Summing Up Elements of a List

A program to compute the sum of elements of a list can be defined as follows.

```
sum([], PSum, Sum) :- Sum = PSum.
sum([One|Rest], PSum, Sum) :-
    NewPSum := PSum + One,
    sum(Rest, NewPSum, Sum).
```

```
sum(List, Sum) :- sum(List, 0, Sum).
```

This program consists of two predicate definitions. Both have the same name `sum` but with different numbers of arguments (sometimes called *arities*). In KL1, in accordance with the tradition of logic programming, predicates with the same name but different arity are considered to be different predicates. Such predicates are often used for auxiliary purposes, as in this program.

The first two clauses with three arguments is the main routine that actually makes the sum. When this predicate is examined individually, its function is to compute the sum of two items: one is the sum of the elements of the list given as the first argument and the other is the value given as the second argument. This second argument is actually a partial sum up to here. The computed total is returned to the third argument.

The first clause reads: if the first argument is an empty list, as no more elements to be added are there, the second argument should be the result. The second clause means: if the list is not empty, having `NewPSum` be the sum of the first list element `One` and the second argument `PSum`, the total of this `NewPSum` and the rest of the elements in `Rest` should give the grand total.

The last clause defines a two-argument predicate, which is the top level. It means: by adding the sum of elements of a list with 0, the sum of the list elements can be obtained.

Note that, the addition in the body of the second clause and the recursive invocation can be executed in parallel. As the guard condition only tests whether the list is empty or not, the result of the addition is not needed in selecting which of the two clauses to execute. Thus, recursions can proceed without waiting for the additions to complete, making many addition goals to be executed concurrently. As one of the arguments of the additions (`PSum`) is the result of the addition one cycle before, the order will be determined among additions in this case.

Concurrency, that is, possibility of parallel execution need not suggest speed-up by actual parallel execution. Because of overhead of communication, doing such a simple operations as integer additions in parallel is usually not beneficial and KL1 implementations will not try that. However, when the simple additions in this example are replaced by more complicated operations, such as *bignum* divisions, actual parallelism may become really beneficial.

As in this example, what is called a `process` in KL1 is not necessarily a sequential process, but may have concurrency within itself. Whether to consider a process being a process totally depends on the viewpoints.

4.1.3 List of Natural Numbers

A program that makes a list of all natural numbers less than a given number can be defined as follows.

```
naturals(N, M, List) :- N >= M | List = [].
naturals(N, M, List) :- N < M |
    List = [N|Rest],
    N1 := N + 1,
    naturals(N1, M, Rest).

naturals(M, List) :- naturals(0, M, List).
```

This program also consists of two predicates, the top level one with two arguments and an auxiliary one with three arguments.

The predicate defined by the first two clauses makes a list of integers beginning with the first argument up to one less than the second argument and returns it to the third argument. The first clause means: if the first argument is not less than the second, return an empty list. The second clause means: if the first argument *is* less than the second, return a list beginning with the first argument; the rest of the list should begin with one more than the original first argument.

The last clause defines the top level predicate. By calling the three argument auxiliary predicate with its first argument 0, the desired list will be obtained.

In this example also, the goals in the second clause, i.e., a body unification, an addition and a recursive invocation are executed concurrently (in any order or in parallel).

4.2 Communication

In this section, ways to combine many processes described in the previous section are discussed, to construct programs for more complicated computation.

4.2.1 Combining Processes

Two or more processes can be combined so as to the result of one process can be used in computation of other processes. For example, sum of natural numbers less than a given integer can be computed by a program combining the two programs explained above, defined as follows.

```
sum_up_to(N, Sum) :- naturals(N, List), sum(List, Sum).
```

In this program, the variable `List` is shared by two processes, serving as the medium to pass the list made by the process `naturals` to the process `sum`.

As the invocations of these two processes are written in a body of the same clause, their execution order is not defined. As the predicate `sum` with two arguments has not guard conditions, it will call three argument `sum`. As this depends its clause selection on the value of its first argument, it will be suspended until its value is determined by another process.

Let us review here the definition of the predicate that creates a list of natural numbers.

```
naturals(N, M, List) :- N >= M | List = [].
naturals(N, M, List) :- N < M |
    List = [N|Rest],
    N1 := N + 1,
    naturals(N, M, Rest).
```

As mentioned above, the goals in the second clause, i.e., a body unification, an addition and a recursive invocation can be done in any order. As the unification can precede other goals, the result can be returned *before* the computation ends. Of course, the result has not been computed completely yet. The cdr of the list is still left undefined and thus the length of the list is not determined yet. This is a typical incomplete data structure. However, the result is already known to be a *cons* by this unification, and its car is defined to be the first argument passed, that is, an integer 0 on its first invocation.

Let us now look again into the summing up predicate, the three-argument `sum`.

```
sum([], PSum, Sum) :- Sum = PSum.
```

```
sum([One|Rest], PSum, Sum) :-
    NewPSum := PSum + One,
    sum(Rest, NewPSum, Sum).
```

As the result of the predicate `naturals` has been decided to be a cons, the guard condition (actually written in the head) can be checked out and the second clause will be selected. The car part of the first argument (here, the variable `One` matches against) is also determined to be 0. As another 0 is passed in the original invocation to the second argument `PSum`, the addition `0+0` is already executable, after only a little step in the `naturals` process.

The recursive invocation cannot proceed like this. As the process `naturals` only made its first step, the variable `Rest` is still undefined. In the recursively called `sum`, the value of the first argument is not defined, making it impossible yet to select a clause. This `Rest` is now shared by two processes, `naturals` and `sum`, the former generating its value and the latter consuming it. The recursive invocation of `sum` cannot proceed before `naturals` makes further steps forward. As `naturals` proceeds each step, `sum` also can proceed one more step.

As in this example, returning an incomplete structure as its value, gradually defining its elements afterwards, is a very convenient programming style of KL1 for higher concurrency. If the result should not be returned before the whole structure is completed, the execution of `sum` will be suspended until the execution of `naturals` completely terminates, lowering the parallelism.

Note again that concurrency does not always mean that actual parallel execution is beneficial. Physical parallelism should be decided considering other issues such as communication overhead.

4.2.2 Stream Communication

The two processes explained in the last section are in relation that values generated incrementally by the process `naturals` are consumed also incrementally by the process `sum`. The two processes shared an incomplete data structure, which is actually a list whose elements are incrementally defined.

This list structure can be regarded as a communication channel between two processes. From this viewpoint, the process `naturals` sends natural numbers one by one, while the other process `sum` receives the numbers one by one to proceed the computation. If the data to be received is not available yet, the computation of the process will suspend.

The communication channel is realized by data structures of type *cons*. Their *car* parts convey the data flowing through the channel and their *cdr* parts represent the rest of the channel. At the end of the communication, that is, when there are no more data to convey, the cdr will become `[]`. As the communication channel is represented as a data structure, the order of data in the channel is defined by the data structure completely. The order of execution can never change the order of data. Such a communication channel is called a *stream*.

Data conveyed by streams can be regarded as a *message* exchanged between processes. In this viewpoint, the process `natural` is sending requests to add numbers to the process `sum`.

4.2.3 Process State

Processes perform computation according to the messages they receive. If what should be done is completely determined by the messages, there will be no need to write in that way. In stead of making a process and sending messages, defining a predicate and calling it would suffice.

The merit of using message-driven processes is that processes can have their *state*. For example, a process which holds an integer value as its state and changes the value according to the received messages can be defined as follows.

```
counter(Stream):- counter(Stream, 0).

counter([], Count).
counter([up|Stream], Count) :-
    NewCount := Count + 1,
    counter(Stream, NewCount).
counter([down|Stream], Count) :-
    NewCount := Count - 1,
    counter(Stream, NewCount).
```

The process defined in this program has the value of the counter as its second argument. Its initial value is 0, as invoked from the top level predicate, which will be incremented and decremented each time `up` and `down`" message arrives. As in this example, a process in KL1 retains its state as its arguments.

To use such a process, the predicate should be called with an argument used for the message stream as follows.

```
?- counter(Stream), some_other_process(Stream).
```

This message stream will be the only interface with this process and the others. The value of the counter retained as the process status can never been accessed directly from outside. Using arguments to retain the process state is encapsulating the information, hiding it from outside.

How such information, then, can be accessed? In the above program, the counter values can be altered by sending messages, but the value can never be read out. It will be explained in the next section.

4.2.4 Structured Messages

Up to here, all messages sent through streams were atomic data. Structured data, of course, can be used as messages.

Functor structures are frequently used as messages. As explained above, functor structures are actually vectors with their first element (the functor name) being a symbolic atom representing the kind of structure. Functors are conveniently used as messages, the functor name indicating the kind of message and other elements (arguments of the functor) describing details.

The above example program of a counter can be revised to allow increment and decrement by arbitrary amount by adding arguments to messages, as follows.

```
counter(Stream):- counter(Stream, 0).
```

```
counter([], Count).
counter([up(N)|Stream], Count) :-
    NewCount := Count + N,
    counter(Stream, NewCount).
counter([down(N)|Stream], Count) :-
    NewCount := Count - N,
    counter(Stream, NewCount).
```

To read out the counter value, the following clause, which also defines functor message, can be added to the program.

```
counter([show(Value)|Stream], Count) :-
    Value = Count,
    counter(Stream, NewCount).
```

The message `show` for reading out the counter value should be sent with the argument `Value` left undefined, to let the counter process return the value there. Such messages with undefined arguments are called an *incomplete message*. The incomplete message mechanism is another example of the use of incomplete data structures. Using this style, a single stream can be used for both-way communication.

4.3 Summary of Processes and Streams

Programming style of KL1 using processes and streams is explained.

- Processes of KL1 are recursively calling goals and their internal states are their arguments. Manipulation of incomplete lists allows concurrent execution of processes.
- Message streams connecting processes are realized by list data structures. Incomplete messages can be used for returning values that depend on internal states back to the message sender.

Incomplete data structures play significant roles in this programming style.

5 Process Network

The programming style of KL1 with processes and message streams connecting them is explained in the previous chapter. In this chapter, various typical methods of combining processes with streams to construct more complex programs are described.

5.1 Filters

A program to compute the sum of natural numbers less than the given value was explained in the previous chapter. The program consisted of two processes, one generating a list of natural numbers and another summing up the elements of the list.

Let us now consider a slightly different problem of computing the sum of *squares* of natural numbers up to the given value. Two methods to make the program revising one in the previous chapter might be as follows.

- Altering the process generating a list of natural numbers so that it will create a list of squares of them.
- Altering the process summing up the list elements so that it will sum up the squares of them.

These two methods are both trying to alter the existing program. Such methods, of course, will work and probably good enough for this particular problem. If, however, both processes were much more complicated and their alteration was much more difficult, could there be any better ways to complete the task?

A program that makes a list of natural numbers and a program that sums them up are already there. If a program that converts a list of integers to a list of the squares of each element can be made and put in between the two processes, it can be an excellent answer. The following program does this.

```
square([], Out) :- Out = [].
square([One|Rest], Out) :-
    Square := One * One,
    Out = [Square|OutTail],
    square(Rest, OutTail).
```

With this predicate defined, the top level program will be as follows.

```
square_sum_up_to(N, Sum) :-
    naturals(N, Naturals),
    square(Naturals, Squares),
    sum(Squares, Sum).
```

The predicate `square` takes a list as its input and creates another list as its output. How this looks like when the this predicate is regarded as implementing a process and two lists as streams? In this view, the program defines a process that receives messages (which are integers) from one stream and send messages (which are square of the received messages) to another stream. A process like this, making some modification to the message received from an input stream and send the modified message to the output stream, is called a *filter*. The overall structure of this program is that processes `naturals` and `sum` are communicating through the filter `square`.

In this particular example, the process `square` playing the role of a filter does not have any internal state, sending messages depending only on the received messages. In general, outgoing messages can depend not only on incoming messages but also on the state of filter processes. As the state of processes may be altered by incoming messages, filtering function can be history sensitive.

Moreover, an input message might not correspond to an output message; a filter may receive multiple messages before sending out one or it may send out multiple messages on receiving a single message. The following program defines a filtering predicate that sums up integer numbers as specified by input messages `add` and outputs the current sum on receiving a message `sum`.

```
accumulate([], Out, Sum) :- Out = [].
accumulate([add(N)|In], Out, Sum) :-
    NewSum := Sum + N,
    accumulate(In, Out, NewSum).
accumulate([sum|In], Out, Sum) :-
    Out = [Sum|NewOut],
    accumulate(In, NewOut, Sum).

accumulate(In, Out) :- accumulate(In, Out, 0).
```

This filter is history sensitive and it does not have one-to-one input/output message correspondence.

5.2 Concatenating Streams

Let us alter the problem slightly again. Now the problem is to compute the total of both squares and cubes of natural numbers less than the given integer.

A similar method as the solution to the previous problem can be used, making a filter which generates `n * n + n * n * n` for input `n`. This, however requires revision on the already working predicate `square` and the revision also makes it more complicated (Again, if this particular example problem is all you need, you might have better rewrite everything from scratch. The method described here is actually useful when required computation is much more complicated and rewriting filters requires considerable efforts). Let us then leave the `square` predicate as it is and define another filter predicate that sends cubes of input message.

Although you now should be quite certain how to define such a predicate. It can be defined as follows.

```
cube([], Out) :- Out = [].
cube([One|Rest], Out) :-
    Cube := One * One * One,
    Out = [Cube|OutTail],
    cube(Rest, OutTail).
```

The outgoing messages of the `naturals` process can be sent to both `square` and `cube` processes by simply letting them read the same message stream as follows.

```
square_sum_up_to(N, Sum) :-
    naturals(N, Naturals),
```

```
square(Naturals, Squares),
cube(Naturals, Cubes),
...
```

Note that the stream `Naturals` here simply represents a data structure. Messages in it will never be lost just because somebody read it. To obtain messages in the stream, processes have to examine through the data structure explicitly. Letting both two processes read all the messages from the same message stream, thus, is quite simple as given above.

A more difficult problem is how the messages in the output streams, `Squares` and `Cubes` in the above program, can be fed to the process `sum`.

One method is to let all the messages from one stream go first, and after all of them are sent, let messages from the other go. A predicate to do this switching can be written as follows.

```
append([], In2, Out) :- Out = In2.
append([Msg|In1], In2, Out) :-
    Out = [Msg|OutTail],
    append(In1, In2, OutTail).
```

This predicate takes two input streams as the first and the second argument, and the third argument will be the output stream. The overall function of this predicate is sending messages from the first input stream to the output stream first, and sending message from the second input after the first stream comes to an end.

When there remain some messages in the first input stream, that is, when the first argument is a not-empty list, the second clause is selected; the first message is sent to output. When the first input stream comes to an end, that is, the first argument becomes an empty list, the first clause is selected; the unification in the body of the first clause connects the output stream directly with the second input stream, which means that all the message coming from that stream should be automatically transferred to the output without receiving and sending them individually.

If this predicate is to be interpreted as a predicate on lists rather than on streams, it makes a concatenation of two lists given as the first and the second arguments and returns it to the third argument.

Using this predicate, the top level of the program will be as follows.

```
queer_sum(N, Sum) :-
    naturals(N, Naturals),
    square(Naturals, Squares),
    cube(Naturals, Cubes),
    append(Squares, Cubes, Both),
    sum(Both, Sum).
```

Note that, in the `append` program, the values of the messages are not examined at all. It only checks whether there are any messages coming or not and forwarding them without looking into them. Thus, the same `append` program can be used for streams with any kind of messages. This benefits from the fact that variables of KL1 do not have specific type declared.

5.3 Merging Streams

The method described in the previous section is not completely satisfactory. In that method, messages from the stream `Cubes` will not arrive the process `sum` until the stream `Squares` is closed. If the process `square` is time consuming, the process `sum` running concurrently with it will become idle. As the process `cube` may have already sent some output to the stream `Cubes`, those might be used by `sum`, but, as the stream is choked at `append`, they cannot reach `sum`.

Let us then consider a method in which, before one input stream coming to an end, messages from the other stream can also be forwarded. Messages coming from either stream should be forwarded to the output. The following program defines such a predicate.

```
merge([], In2, Out) :- Out = In2.
merge(In1, [], Out) :- Out = In1.
merge([Msg|In1], In2, Out) :-
    Out = [Msg|OutTail],
    merge(In1, In2, OutTail).
merge(In1, [Msg|In2], Out) :-
    Out = [Msg|OutTail],
    merge(In1, In2, OutTail).
```

The first two clauses mean when either of the input streams comes to an end, the other stream should be directly connected to the output. The remaining two clauses are for forwarding messages from either of the input streams to the output.

A process forwarding multiple input streams to an output stream is called a *merger*. The `append` predicate described above is also a kind of merger, but with stronger restriction.

Examine the program carefully to see what will happen if both streams already have messages to be forwarded. In such a case, both the third and the fourth clauses have their clause application condition satisfied. As stated many times before, the language does not define which clause will be selected. The same merger program with the same input data thus can send messages to the output stream in different order. This is a typical example where the *nondeterminacy* of KL1 is made apparent and this nondeterminacy sometimes makes debugging more difficult. It is also possible, of course, to write deterministic programs as `append` to avoid the problem, but nondeterminacy can be indispensable for gaining concurrency as in this example. Note that, although the order of messages from different input streams is nondeterministic, the order between messages originally from the same input stream will be preserved in the output.

Using this `merge`, the top level program will be as follows.

```
queer_sum(N, Sum) :-
    naturals(N, Naturals),
    square(Naturals, Squares),
    cube(Naturals, Cubes),
    merge(Squares, Cubes, Both),
    sum(Both, Sum).
```

With this program, messages from both processes `square` and `cube` can be forwarded to the process `sum` as soon as they become available, giving the program better concurrency than using `append`.

5.4 Builtin Merger

The merger defined by user in KL1 explained in the previous section still has problems. One problem is with the number of input streams. As predicates with varying number of arguments cannot be defined in KL1, only mergers with fixed number of input streams can be defined. Thus, to realize mergers with dynamically changing number of inputs, a tree-like network of mergers has to be made, which allows addition of new input streams dynamically. When a message arrives at one of the input streams of such a network, it must go through as many mergers as the height of the tree structure. An unbalanced tree structure, thus, may result in very low performance. The tree structures can be balanced on addition and deletion of input streams by any of the well-known tree balancing algorithms, but this makes the program more complicated.

Even if the tree is completely balanced, a merger network with n input streams will have height proportional to $log\ n$ and each message has to go through $log\ n$ mergers. On the other hand, in procedural concurrent programming languages, using locking and destructive assignment, it is quite easy to realize a merger with a constant overhead. As mergers are very frequently used in process and stream programming style of KL1, difference in order of computation is unbearable.

To solve this problem, KL1 provides mergers as one of its builtin features. Built-in predicate for stream merging has two arguments and is invoked as `merge(In, Out)`. When a message is sent to the input stream (that is, when the first argument is given its value a cons structure with the message as its car), it is simply forwarded to the output stream (the second argument is unified with another cons structure having the message as its car). The process starts working as a merger when, instead of a cons, a vector structure such as `{ In_1, In_2, ..., In_n }` is given as its first argument. The merger then becomes an n input merger. The number of inputs can be increased by giving any one of the input `In_k` a vector value again. To decrease the number, simply closing input streams (that is, giving them a symbolic atom `[]` as their value) will suffice. When all the inputs are closed, the output is also closed (unified with `[]`).

The built-in merger can forward messages with constant overhead irrespective of the number of input streams, and the constant is also much smaller than mergers defined in KL1.

5.5 Dispatchers

In the above program computing the total of both squares and cubes, both of the filters making squares and cubes required same inputs. If that is not the case, that is, if some of the messages should go to different processes depending its contents, how should it be programmed?

Let us, this time, consider a still more meaningless example, in which we compute the sum of squares of even numbers and cubes of odd numbers for natural numbers less than the given value. To do this, we can use a process which dispatches the numbers to one of two filters depending on whether the input is even or odd. Such a program can be defined as follows.

```
dispatch([], Odds, Evens) :- Odds = [], Evens = [].
dispatch([One|Rest], Odds, Evens) :- One/2 =\= 0 |
    Odds = [One|OddsTail],
```

```
        dispatch(Rest, OddsTail, Evens).
   dispatch([One|Rest], Odds, Evens) :- One/2 =:= 0 |
        Evens = [One|EvensTail],
        dispatch(Rest, Odds, EvensTail).
```

The predicate defined by this program has one input stream and two output streams. The received messages are sent to either of the output streams depending on their contents. Such a process is called a *dispatcher*.

Using this predicate, the top level will be as follows.

```
   queer_sum(N, Sum) :-
        naturals(N, Naturals),
        dispatch(Naturals, Odds, Evens),
        square(Evens, Squares),
        cube(Odds, Cubes),
        merge(Squares, Cubes, Both),
        sum(Both, Sum).
```

5.6 Servers

Another important technique for constructing process networks is using processes called a *server*. Servers have an input stream which usually is connected to many *client* processes through a merger. Client processes send data to be shared among clients in messages to the server and the server keeps them, in some suitable format, as its state (actually, arguments of its recursive invocation). Clients can also read out the data by sending incomplete messages for query.

The process of `counter` explained in the previous chapter is a typical server process.

5.7 Summary of Process Network

KL1 programs are often constructed with many processes connected by streams. Typical element processes of such process networks and the ways to connect them were explained in this chapter.

- Filters have one input and one output stream. They receive input messages and send out output messages either depending or not depending on their internal states, which may depend on messages received so far.

- Two or more streams can be combined into one by either concatenation or merging. KL1 provides builtin mergers that efficiently merges many streams into one.

- Dispatchers have multiple output streams that forwards input messages to some of the output streams depending on their contents.

- Servers are connected to many clients through merged streams. Incomplete messages can be used for communication between servers and clients.

6 Pragmas

This chapter explains about *pragma* which specifies how concurrent programs that *can run* in parallel *should actually run* in parallel.

6.1 Principles of Specification of Parallel Execution

Before going into details, the principles made in the design of KL1 on how to specify parallel execution are described in this section.

6.1.1 Programs Specify Load Distribution

To perform computation efficiently on actual parallel processing hardware, the following two aspects must be considered.

Distribution of Load

> If computation is concentrated on one or small number of processors, the whole computation will not end until those processors finish the job, while making all other processors idle waiting for them. Thus, the problem to be solved *should* be divided into subproblems small enough and distributed to as many processors as possible, giving equal load to all the processors.

Reduction of Communication:

> To distribute subproblems, communication is needed for distributing such subproblems and also for gathering the results. Also, in most efficient algorithms, subproblems are not completely independent, requiring communication between processors solving different part of the problem during the computation. If too much communication is required, its cost may dominate the computation cost, sometimes making parallel execution less efficient than sequential execution. Thus, the problem should *not* be divided into too small pices in order to reduce the amount of communication between the solvers.

These two targets of load distribution and communication reduction conflict each other: too much distribution may result in too much communication; too much reduction of communication won't allow load distribution. Thus, finding a method of division into subproblems without too much communication and finding a reasonable trade-off point of communication cost and load imbalance cost is a central issue of research and development of parallel processing software.

How a problem can be divided into subproblems depends heavily on the problem itself and algorithms used. When the amount of computation can be predicted accurately, as in computing products of many matrices, automatic load distribution might be able to do a decent job. In the area of knowledge information processing, however, the result of solving one subproblem might drastically change the amount of computation needed to solve other subproblems, making automatic load distribution very difficult with the current level of parallel processing software technology.

As this is the status quo, KL1 is designed as a *tool* for the research automatic load distribution. Load distribution is thus not automatic by the language implementation, but specified in the programs. So as to ease experimentation of diverse load distribution methods, the language is designed to make changes in load distribution easy, which will be discussed in more detail in the next section.

6.1.2 Separating Concurrency and Parallelism

In KL1, *concurrency* and *parallelism* are separately specified. Here, we use the two words in the following meanings.

Concurrency

> specifies which part of the program *can* be executed in parallel. Concurrency affects the correctness of the programs. If program parts that should never be executed in parallel are made concurrent, the program might not work correctly. By specifying clause application condition in the guard, clause selection will be automatically suspended until data required for checking the condition become available. This *data-flow synchronization* implicitly specifies allowed concurrency.

Parallelism

> specifies which concurrency should be actually exploited for efficient parallel execution. This is specified by *pragma* in programs. As pragma only specifies which of concurrently executable portions of computation should really be executed in parallel, it will never affect the correctness of the program. However, it may affect the *efficiency* of the program drastically.

During research of efficient load distribution methods, diverse methods have to be tested. If debugging of parallel programs, which, even done only once, already demands more effort than for sequential programs, has to be done for each time a new method is tested, the required effort might become enormous. By separating concurrency and parallelism, KL1 allows load distribution methods to be altered without affecting the correctness of programs.

6.2 Goal Distribution

For distribute computation, the *node* to execute a goal is specified. A node is a processor or a group of processors within which load distribution is automatic. As communication cost is relatively small among processors physically sharing a bus or memory, automatic load distribution may be tried among such group of processors. For large scale systems, however, such tight coupling may not be appropriate, making communication cost much higher.

Execution node specifications can be given to body goals with the following notation.

```
Goal @ node(Node)
```

In the current implementation, nodes are specified by their sequential numbers starting from zero and up to the number of available processors minus one.

An example of a program specifying goal distribution follows.

```
p([One|Rest], N, State) :-
    q(One, State, NewState)@node(N),
    N1 := N + 1,
    p(Rest, N1, NewState).
```

In this example, each message received by `p` are handled on different node by the predicate `q`. Goals without any node specification pragma (an addition and a recursive invocation of `p` in this example) are executed in the same node as the original goal.

6.3 Goal Priority

When there are more concurrently executable goals than the number of physical processors, the order of their execution sometimes affects the efficiency drastically. Thus, to suggest which goal should be executed earlier for better efficiency, goals can be given a priority by a pragma.

With the KLIC implementation, priority specifications can be given to body goals with one of the following notations.

```
Goal @ priority(AbsPrio)
Goal @ lower_priority(RelPrio)
Goal @ lower_priority
```

Here, *AbsPrio* and *RelPrio* should be a non-negative integer constant, or a variable which should be instantiated to a non-negative integer later. In the current implementation, negative priority values are interpreted as zero.

With the absolute priority specification, the goal with the specification will have the priority value specified. With the relative priority specification, the goal will have priority less than the priority of the parent goal by the specified amount. The specification `Goal@lower_priority` has the same effect as `Goal@lower_priority(1)`. Goals without any priority specifications will have the same priority as their parents.

The highest possible priority is the largest possible integer value, which depends on host systems (see Section 3.2.2 [Numerical Atoms], page 10). The initial goal `main:main` has the maximum priority possible for the host system.

Priority pragmas are not necessarily obeyed faithfully, as it may prevent efficient implementation. As pragmas do not affect the correctness of programs, the implementation may safely neglect some pragmas. In the current implementation, priority specification is effective only among goals in the same node.

6.4 Clause Priority

When more than one clause have their guard condition satisfied, the language does not define which will be selected. Programs should be written so that any selection will yield the correct result. However, even if the program is written so that selection of clauses will not affect the correctness, it does sometimes affect efficiency of programs considerably. Priority between clauses can be specified by a pragma to suggest which clause should be selected for better efficiency.

To specify preference among clauses of a predicate, write clauses with higher priority first, a pragma `alternatively` followed by a period, and then the rest of clauses which have lower priority. The `alternatively` pragma can appear more than once. No priority relation is assumed within clauses before the first `alternatively`, after the last one, or between two of them.

The following program defines a merge process that forwards messages from the stream given as the first argument in prior to the messages from the second stream when both are available.

```
merge([], In2, Out) :- Out = In2.
merge(In1, [], Out) :- Out = In1.
merge([Msg|In1], In2, Out) :-
```

```
        Out = [Msg|OutTail],
        merge(In1, In2, OutTail).
    alternatively.
    merge(In1, [Msg|In2], Out) :-
        Out = [Msg|OutTail],
        merge(In1, In2, OutTail).
```

Note however that if the second input stream has a message and the first does not, it will forward one from the second.

Clause priority specification, also, might be neglected sometimes, for more efficient implementation.

6.5 Summary of Pragmas

This chapter described the *pragma* feature of KL1, which suggested the language implementation the way to realize more efficient processing.

- KL1 is designed with the following principles.
  - Load distribution is specified by programs.
  - Logical concurrency and physical parallelism are specified separately.
- Goal distribution is specified by `@node` pragma.
- Goal priority is specified by priority pragmas.
- Clause priority is specified by `alternatively` pragama.

It may be worthwhile to note here again that pragmas are only suggestions for efficiency, not specifying how the implementation should behave rigidly. Pragmas may be neglected sometimes. Thus, programs cannot depend on pragmas to be correct.

Concept Index

U

V